

Predicting Duke vs. UNC Feb. 3, 2023 using modeling

Dom Fenoglio, The Duke Chronicle


What is the game prediction?

	date	team	opp	game_location	tempo	ppp	pts	win_per
1	2024-02-03	Duke	North Carolina	A	71.05691	1.037392	73.7	22.89543
	did_win	simulate_date	year					
1	FALSE	2024-02-03	2024					

	date	team	opp	game_location	tempo	ppp	pts	win_per
1	2024-02-03	North Carolina	Duke	H	71.05691	1.152916	81.9	77.10457
	did_win	simulate_date	year					
1	TRUE	2024-02-03	2024					

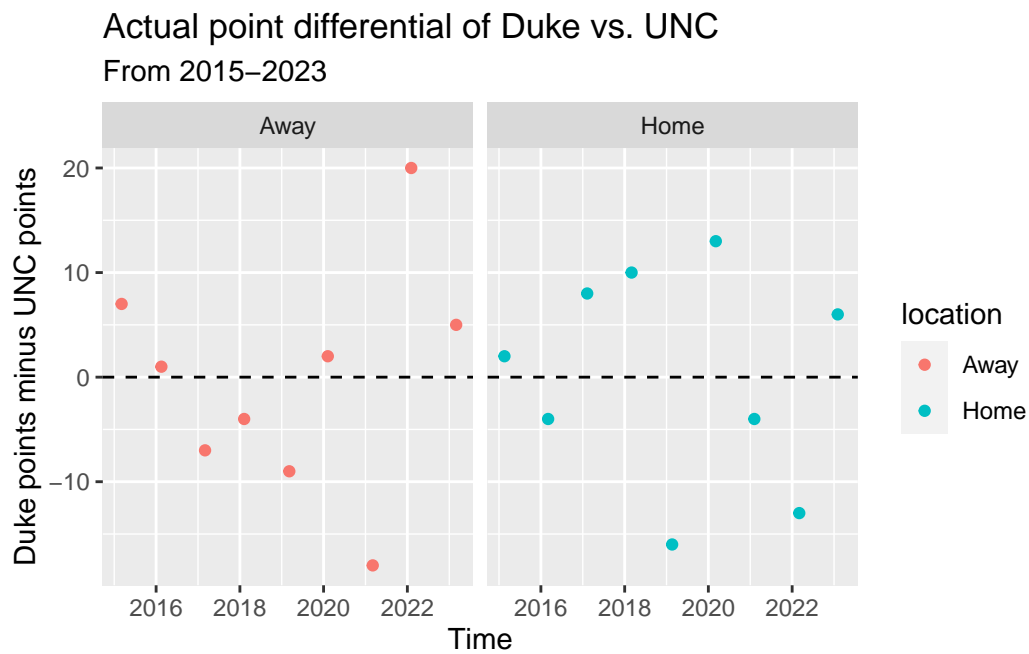

100 simulations of Duke vs. UNC

Using probability 0.22895 for a Duke win



This is a simple representation of 100 simulations of a binomial trial with the given probability of Duke winning. In other words, if the game was played 100 times, these would be the results. Clearly, the Tar Heels are going to run away with this one. But, this is just a surface level look at the data, so let's look closer.

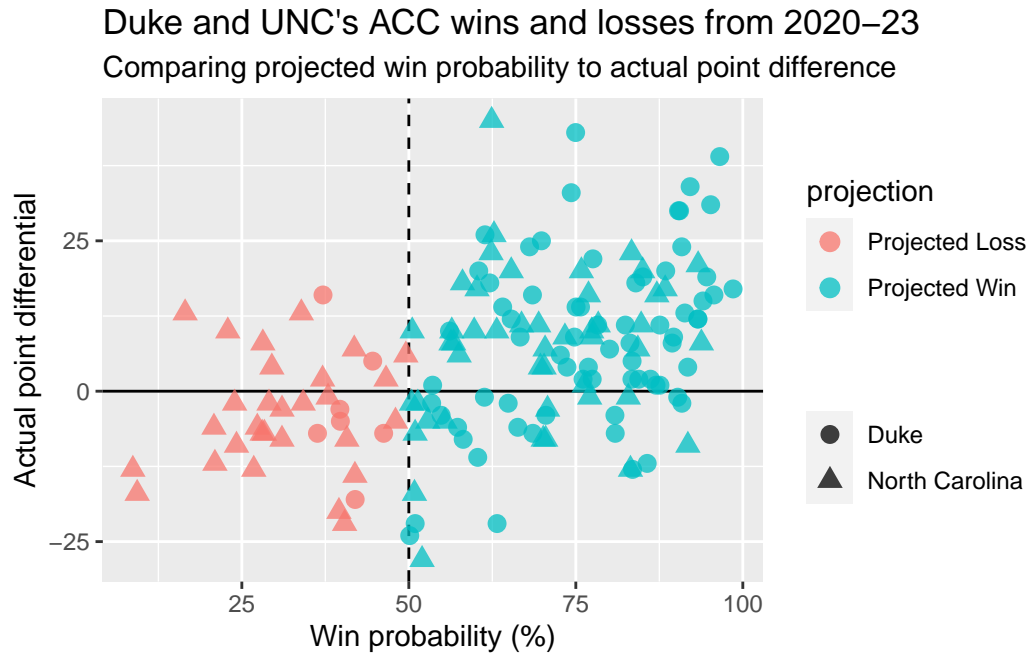
Some exploratory analysis



These graphs allow you to see the vast difference between projected point differentials and actual game scores. In particular, notice the blue dot dated right after 2022. This was Coach

K's final home game. An unranked North Carolina team beat No. 9 Duke in a decisive 13-point victory. On the whole, the predictions can get really far off. Maybe it's just hard to predict basketball games (it is), and the model can't do much.

Duke and UNC ACC games 2020-2023



Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	25	11
1	33	85

Accuracy : 0.7143
 95% CI : (0.636, 0.7841)
 No Information Rate : 0.6234
 P-Value [Acc > NIR] : 0.011316

 Kappa : 0.3421


 McNemar's Test P-Value : 0.001546

Sensitivity : 0.8854
 Specificity : 0.4310
 Pos Pred Value : 0.7203
 Neg Pred Value : 0.6944
 Prevalence : 0.6234
 Detection Rate : 0.5519
 Detection Prevalence : 0.7662
 Balanced Accuracy : 0.6582

'Positive' Class : 1

This graph allows us to see the distribution of Duke and North Carolina's ACC games from the past few years, and how the model did at predicting the outcome. On the whole, it did a pretty good job, as shown by the confusion matrix analysis. Great! Then, it should be able to be just effective at predicting Duke vs. UNC, right?

Now, let's see the past 8 years of Duke vs. UNC



This graph breaks down win probability versus actual results, with coloring showing which shade of blue won. A quick look at this graph shows that our model didn't do a very great job. Why are these games so far off if on the whole it does a good job? Again, let's look closer.

Confusion Matrix and Statistics

```

      Reference
Prediction 0 1
0 3 3
1 5 7

      Accuracy : 0.5556
      95% CI : (0.3076, 0.7847)
No Information Rate : 0.5556
P-Value [Acc > NIR] : 0.5966

      Kappa : 0.0769

McNemar's Test P-Value : 0.7237

      Sensitivity : 0.7000
      Specificity : 0.3750
Pos Pred Value : 0.5833
Neg Pred Value : 0.5000
Prevalence : 0.5556
Detection Rate : 0.3889
Detection Prevalence : 0.6667
Balanced Accuracy : 0.5375

      'Positive' Class : 1
```

In fact, this data does not tell us much about the true outcome of the game at all. Positive and negative predictive values are low, and we might be better off just flipping a coin. While we can't create any concrete reasons why solely from this data, we can conclude that this model does not consistently provide an accurate prediction.

Conclusions

In my own humble opinion, this game is the best rivalry in sports. Every piece of data can be thrown out the window when it comes to these two schools, and instead of trying to jump to

any conclusions, we're better off just going along for the ride. My full article breaking down these findings in a more compact, thorough way can be found on dukechronicle.com.